Homework 27

Brian Knotten, Brett Schreiber, Brian Falkenstein

October 30, 2018

1

1.1

A parallel algorithm for AND with $p = n$ is given:

```
P_And(x_1 ... x_n, p):
   if p == 1:
       return and(x_1...x_n)
   else:
       return and(P_And(x_1...x_(n/2), p/2), P_And(x_(n/2+1) ... x_n, p/2))
```

In the case where $p = n$, the algorithm will cascade down until each processor has a single x value (n values, n processors, each gets one). After this step, the ands will cascade upwards, with one processor anding the result of 2 processors. Thus, at the first time stamp, all n processors are utilized (although, they do not really do anything, they just return the and of the single value they have), at the second time stamp, $n/2$ processors are utilized, at the third time $n/4$ processors, etc.

Note that this algorithm is EREW. It is exclusive read because no two processors are reading the same value at the same time. The processors simply pass their answers on to another processor, which then further processes the answers. It also only writes one time, which is the processor at the end that returns the and of the entire input. The efficiency is clearly bad. Given the equation for efficiency:

$$E(n,p) = \frac{S(n)}{pT(n,p)}$$

Clearly, $S(n)$ will be n, as it takes n time to and an input of size n (just and the first two, then the result of that with the third, etc.). We can further define $T(n,p)$, the recurrence relation, as:

$$T(n,p) = T(n/2, p/2) + 1 = log(n)$$

This recurrence relation is true because at each step, we are simply doing the and of 2 elements, which takes constant time. Further, these ands happen simultaneously. Thus, we can compute the efficiency as:

$$E(n,p) = \frac{S(n)}{pT(n,p)} = \frac{n}{nlog(n)} = \frac{1}{log(n)}$$

For large n values, this efficiency is rather bad.
Using the folding principle, which states:

$$T(n,p) \leq kT(n,kp)$$

We can see that if we have $p = n^{1/3}$ processors instead of n, we'd get an upper bound on the running time of:

$$T(n,p) \leq \frac{1}{n^{2/3}}T(n, n^{1/3})$$

IE, reducing the number of processors from n to $n^{1/3}$, at most increases the run time by $1/n^{2/3}$.

1.2

The parallel algorithm for AND with $p = n/log(n)$ is the same as the algorithm when $p = n$. The difference here is that at the base level, where in the last algorithm each processor had 1 value to pass up, each processor will have $log(n)$ values that it must sequentially and. This is because the problem is split into equal size sub problems, in this case n/p, IE each processor initially gets a subproblem of size n/p. In the case where $p = n$, $n/p = 1$. However, when $p = n/log(n)$, plugging in, we get $n/p = log(n)$. That means, instead of the first step of the algorithm taking constant time, it now takes $log(n)$, as all the processors sequentially and their $log(n)$ values. Beyond this, the algorithms function almost identically, however this case will spend less time cascading answers up than the one with $p = n$. This leads us to a recurrence relation of:

$$T(n, p) = n/p + log(p) = log(n)$$

And an efficiency of:

$$E(n, p) = \frac{S(n)}{pT(n,p)} = 1$$

Again using the folding principle, we can set an upper bound on the run time for this algorithm if we reduce the number of processors from $p = n/log(n)$ to $p = n^{1/3}$.

$$\frac{n}{log(n)}k = n^{1/3}$$

$$k = \frac{n^{1/3}log(n)}{n}$$

Thus, our time increases by at most $\frac{n^{1/3}log(n)}{n}$ if we decrease the number of processors from $p = n/log(n)$ to $p = n^{1/3}$.

1.3

Without the restriction of an EREW machine, and with n processors, we can find an algorithm for AND that takes constant time. In the first time step, the algorithm copies 1 of the input data into each processor. Then, each processor executes the following code:

```
processorAnd(x1):
    if x1 == 0:
        ans = 0
```

Where ans is some variable that all processors have write access to, and is initialized to be 1. Thus, in the second time stamp, each processor checks their value, and either does nothing if they have a 1, or they write a 0 if they have a 0. The solution will then be stored in ans, which will be 1 if all the inputs are 1 (no processor overwrites it with a 0), or 0 if even one if the inputs is 0. Because we simply have a constant number of operations that each take constant time, this algorithm is O(1).

Here, since there is no recurrence, our runtime T is $T = 1$. Then, our efficiency is:

$$E(n, p) = \frac{S(n)}{pT(n,p)} = \frac{n}{n * 1} = 1$$

If we decrease the number of processors from n to $n^{2/3}$, by the folding principle, we can get an upper bound on the run time:

$$T(n, p) \leq kT(n, kp)$$

$$T(n, p) \leq \frac{1}{n^{1/3}}T(n, n^{1/3})$$

3

3.1

The EREW algorithm with $p = n$ is defined:

```
createArray(n, x, p):
   if p == 1:
      A = new Array[n]
      for i in range(0, len(A)):
            A[i] = x
      return A;
   else:
      return concat(createArray(n/2, x, p/2), createArray(n/2, x, p/2))
```

That is, with $p = n$, at the first time stamp, each processor creates an 'array' of size 1 and sets its value at index 0 to be x. Then, cascading up the call tree, each processor concatenates the arrays from 2 processors, creating a larger array. So at time 2, some processor will combine 2 size 1 arrays to create a size 2 array. This continues until the final array of size n is constructed. The first step of the algorithm takes constant time, as each processor simply creates a new array object and stores a value in it (and they all do this concurrently). It then takes $log(n)$ time to cascade upwards, as at each time step, the processors cut the number of array concatenations left to do in half. This leads us to a recurrence relation of:

$$T(n, p) = T(n/2, p/2) + 1 = log(n)$$

As at each step, the algorithm cuts the number of arrays left to concatenate in half, and each step takes constant time. Note that S(n) = n, as it would take n time to create a size n array and write a value to each index (just do it sequentially). Thus, with $p = n$, we get efficiency:

$$E(n, p) = \frac{S(n)}{pT(n, p)} = \frac{n}{nlog(n)} = \frac{1}{log(n)}$$

Similarly to problem 1, if we decrease the number of processors from n to $n^{1/3}$, by the folding principle, we get an upper bound of:

$$T(n, p) \leq \frac{1}{n^{2/3}} T(n, n^{1/3})$$

3.2

The algorithm here looks exactly the same as the one specified in 3.1 The difference is that whereas that algorithm spent constant time at the leaves, and $log(n)$ time cascading up, this algorithm takes $n/log(n)$ time at the leaves and $n/log(n)$ time cascading. This gives us a time complexity of:

$$T(n, p) = n/p + log(p) = \frac{n}{n/log(n)} + log(n/log(n)) = log(n) + log(n/log(n)) = log(n)$$

And an efficiency of:

$$E(n, p) = \frac{S(n)}{pT(n, p)} = \frac{n}{n/log(n) * log(n)} = 1$$

Note that if we decrease the number of processors from $n/log(n)$ to $n^{1/3}$, we can use the folding principle to get an upper bound on our run time:

$$T(n, p) \leq \frac{1}{n^{2/3}} T(n, n^{1/3})$$

As we are decreasing our number of processors by $k = \frac{1}{n^{2/3}}$.

3.3

Given n processors, and allowing the processors to write and read to the same location in memory, we get an algorithm for process i:

```
processor_i_createArray(x, i):
      A[i] = x
```

Where A is the result array. We can declare A in constant time before having all the n processors write to it. Note that the processors use concurrent read here, as they all read the x variable and write it to the array. However, there is no concurrent writing necessarily, as each processor writes to a different index in the array. This algorithm

is constant time, as it takes constant time to create the array (just define a starting index in memory), and constant time for each processor to concurrently write to the array, and constant time to return the pointer to the array. Because our time $T(n, p) = 1$, we can an efficiency of 1:

$$E(n, p) = \frac{S(n)}{pT(n, p)} = \frac{n}{n * 1} = 1$$

If we decrease the number of processors from n to $n^{2/3}$, by the folding principle, we can get an upper bound on the run time:

$$T(n, p) \leq kT(n, kp)$$

$$T(n, p) \leq \frac{1}{n^{1/3}} T(n, n^{1/3})$$

6

6.1

In a matrix multiplication $A \times B = C$, every entry of C is dependent on a row of A and a column of B, but is independent from every other entry in C. Therefore all of the entries of C can be computed independently at the same time.

With n^2 processors (assuming A and B are both $n \times n$ square matrices) C can be computed in n time with the following CREW algorithm:

MatrixMultiplicationAlgorithm(A, B):
    for each processor $p_{i,j} \in (n \times n)$, concurrently:
        let $C_{i,j} = 0$
        for $k = 1$ to n do:
            update $C_{i,j} += A_{i,k} * B_{k,j}$
    Output C

$T(n, p = n)$ This algorithm takes n time because it makes n multiplication and addition steps at the same time. It is an Exclusive Write algorithm because each processor only writes to one, distinct cell of C. It is a Concurrent Read algorithm because processors assigned cells in the same row in C will both read from the same row in A, and processors assigned cells in the same column in C will both read from the same column in B.

The efficiency of this algorithm is: $E(n, p = n^2) = \frac{S(n)}{p*T(n,p=n^2)} = \frac{n^3}{n^2*n} = 1$.

By the folding principle, the upper bound of the running time for this algorithm on $n^{1/4}$ processors is:
$T(n, p = n^2) \leq \frac{1}{n^{7/4}} \cdot T(n, p = n^{1/4})$

6.2

With n^3 processors (assuming A and B are both $n \times n$ square matrices) C can be computed on a CREW PRAM similarly to the algorithm in 6.1:
Each of the n^2 entries of C will use n processors to compute the n products $A_{i,k} * B_{k,j}$ required for the entry $C_{i,j}$ concurrently in constant time. Then, for each entry, the n products are summed in $log(n)$ time, as at each time step the number of sums being calculated is halved.

$T(n, p = n^3)$ This algorithm takes $log(n)$ time because it takes $log(n)$ steps to sum all n products of each cell of C. It is an Exclusive Write algorithm because each processor only writes to one, distinct cell of C. It is a Concurrent Read algorithm because processors assigned cells in the same row in C will both read from the same row in A, and processors assigned cells in the same column in C will both read from the same column in B.

The efficiency of this algorithm is: $E(n, p = n^3) = \frac{S(n)}{p*T(n,p=n^3)} = \frac{n^3}{n^3*log(n)} = \frac{1}{log(n)}$

By the folding principle, the upper bound of the running time for this algorithm on $n^{1/4}$ processors is:
$T(n, p = n^3) \leq \frac{1}{n^{11/4}} \cdot T(n, p = n^{1/4})$

6.3

With $\frac{n^3}{log(n)}$ processors (assuming A and B are both $n \times n$ square matrices) C can be computed on a CREW PRAM as follows, similar to the algorithms in 6.2 and 6.1:

Each of the processors assigned to each of the n^2 $C_{i,j}$ entries will compute a sub-sum of the required product sums of size $logn$ i.e. each cell in C has $\frac{n}{log(n)}$ processors each computing $\sum_{k=\ell}^{m} A_{i,k} * B_{k,j}$, where ℓ to m is a subset of 1 to n of size $log(n)$. Each of these sub-sums are then summed in $log(n)$ time, as (similar to 6.2) at each time step the number of sums being calculated is halved.

$T(n, p = \frac{n^3}{log(n)})$ This algorithm takes $log(n)$ time because it takes $log(n)$ steps to sum all n products of each cell of C. It is an Exclusive Write algorithm because each processor only writes to one, distinct cell of C. It is a Concurrent Read algorithm because processors assigned cells in the same row in C will both read from the same row in A, and processors assigned cells in the same column in C will both read from the same column in B.

The efficiency of this algorithm is: $E(n, p = \frac{n^3}{log(n)}) = \frac{S(n)}{p*T(n,p=\frac{n^3}{log(n)})} = \frac{n^3}{\frac{n^3}{log(n)}*log(n)} = \frac{n^3}{n^3} = 1.$

By the folding principle, the upper bound of the running time for this algorithm on $n^(1/4)$ processors is:
$T(n, p = \frac{n^3}{log(n)}) \leq \frac{log(n)}{n^{11/4}} \cdot T(n, p = n^{1/4})$

6.4

With $\frac{n^3}{log(n)}$ processors (assuming A and B are both $n \times n$ square matrices) C can be computed on a EREW PRAM as follows:

We will use the same algorithm as in 6.3, however we will first use the $\frac{n^3}{log(n)}$ processors assigned to each cell in C to make n copies of each entry in A and B used to calculate the cell entries in $log(n)$ time. We then apply the algorithm described in 6.3, but the processors assigned to each cell in C now have different locations in memory to access when they access the required elements of A and B.

$T(n, p = \frac{n^3}{log(n)})$ This algorithm takes $log(n)$ time because copying the elements from A and B takes $log(n)$ time and applying the algorithm from 6.3 also takes $log(n)$ time. It is an Exclusive Write algorithm because each processor only writes to one, distinct cell of C. It is an Exclusive Read algorithm because the processors assigned to each cell each have their own unique memory location to access for every element in A and B.

The efficiency of this algorithm is the same as the efficiency of 6.3's algorithm: 1.

By the folding principle. the upper bound of the running time for this algorithm is the same as the upper bound of 6.3's algorithm:
$T(n, p = \frac{n^3}{log(n)}) \leq \frac{log(n)}{n^{11/4}} \cdot T(n, p = n^{1/4})$